

---

**Name:** Jann Moises Nyll B. De los Reyes

**Section:** CPE22S3

**Date:** March 11, 2024

**Submitted to:** Engr. Roman M. Richard

---

## ✓ Linear Regression



In Machine Learning and this notebook we use Scikit-learn a lot

---

### What is scikit-learn used for ?

Scikit-learn (Sklearn) is the most useful and robust library for machine learning in Python. It provides a selection of efficient tools for machine learning and statistical modelling including classification, regression, clustering and dimensionality reduction via a consistent interface in Python.

### What is linear regression used for ?

Linear regression analysis is used to predict value of a variable based on the value of another variable. The variable you want to predict is called *dependent variable*. The variable you are using to predict the other variable's value is called the *independent variable*.

---

## ✓ Making Prediction with Linear Regression

Given the representation is a linear equation, making prediction is as simple as solving the equation for a specific set of inputs.

Let's make this concrete with an example. Imagine we are predicting weight (y) from height (x). Our linear regression model representation for this problem would be:

$$y = B_0 + B_1 * x_1$$

or

$$\text{weight} = B_0 + B_1 * \text{height}$$


Where  $B_0$  is the *bias coefficient* and  $B_1$  is the coefficient for the height column. We use learning technique to find a good set of coefficient values. Once found, we can plug in different height values to predict the weight

For example, let's use  $B_0 = 0.1$  and  $B_1 = 0.5$ . Let's plug them in and calculate the weight (in kilograms) for a person with the height of 182 centimeters.

$$\text{weight} = 0.1 + 0.5 * 182$$

$$\text{weight} = 91.1$$

You see that the above equation could be plotted as a line in two-dimensions. The  $B_0$  is our starting point regardless of what height we have. We can run through a bunch of heights from 100 to 250 centimeters and plug them to the equation and get weight values, creating our line.



Now that we know how to make predictions given a learned linear regression model, let's look at some rules of thumb for preparing our data to make the most of this type of model.

## 📁 Import & Install Libraries

```
1 !pip install hvplot
```

```
Requirement already satisfied: hvplot in /usr/local/lib/python3.10/dist-packages (0.9.2)
Requirement already satisfied: bokeh>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (3.4.1)
Requirement already satisfied: colorcet>=2 in /usr/local/lib/python3.10/dist-packages (from hvplot) (3.1.0)
Requirement already satisfied: holoviews>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (1.17.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from hvplot) (2.0.3)
Requirement already satisfied: numpy>=1.15 in /usr/local/lib/python3.10/dist-packages (from hvplot) (1.25.2)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from hvplot) (24.0)
Requirement already satisfied: panel>=0.11.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (1.4.2)
Requirement already satisfied: param<3.0,>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from hvplot) (2.1.0)
Requirement already satisfied: Jinja2>=2.9 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (3.1.3)
Requirement already satisfied: contourpy>=1.2 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (1.2.1)
Requirement already satisfied: pillow>=7.1.0 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (9.4.0)
Requirement already satisfied: PyYAML>=3.10 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (6.0.1)
Requirement already satisfied: tornado>=6.2 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (6.3.3)
Requirement already satisfied: xyzservices>=2021.09.1 in /usr/local/lib/python3.10/dist-packages (from bokeh>=1.0.0->hvplot) (2024.4.0)
Requirement already satisfied: pyviz-comms>=0.7.4 in /usr/local/lib/python3.10/dist-packages (from holoviews>=1.11.0->hvplot) (3.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->hvplot) (2024.1)
Requirement already satisfied: markdown in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (3.6)
Requirement already satisfied: markdown-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (3.0.0)
Requirement already satisfied: linkify-it-py in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (2.0.3)
Requirement already satisfied: mdit-py-plugins in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (0.4.0)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (2.31.0)
Requirement already satisfied: tqdm>=4.48.0 in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (4.66.2)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (6.1.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from panel>=0.11.0->hvplot) (4.11.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2>=2.9->bokeh>=1.0.0->hvplot) (2.1.5)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->hvplot) (1.16.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->panel>=0.11.0->hvplot) (0.5.1)
Requirement already satisfied: uc-micro-py in /usr/local/lib/python3.10/dist-packages (from linkify-it-py->panel>=0.11.0->hvplot) (1.0.3)
Requirement already satisfied: mdurl>=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py->panel>=0.11.0->hvplot) (0.1.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->panel>=0.11.0->hvplot) (2024.2.2)
```

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 import hvplot.pandas
7
8 from sklearn.model_selection import train_test_split
9
10 from sklearn import metrics
11
12 from sklearn.linear_model import LinearRegression
13
14 %matplotlib inline
```

## 📄 Check out the Data

```
1 df = pd.read_csv('/content/drive/MyDrive/Module 11/Real estate.csv')

1 df.head()
```

	No	X1 transaction date	X2 house age	X3 distance to the nearest MRT station	X4 number of convenience stores	X5 latitude	X6 longitude	Y house price of unit area
0	1	2012.917	32.0	84.87882	10	24.98298	121.54024	37.9
1	2	2012.917	19.5	306.59470	9	24.98034	121.53951	42.2
2	3	2013.583	13.3	561.98450	5	24.98746	121.54391	47.3
3	4	2013.500	13.3	561.98450	5	24.98746	121.54391	54.8
4	5	2012.833	5.0	390.56840	5	24.97937	121.54245	43.1

Next steps: [View recommended plots](#)

1 df.shape

(414, 8)

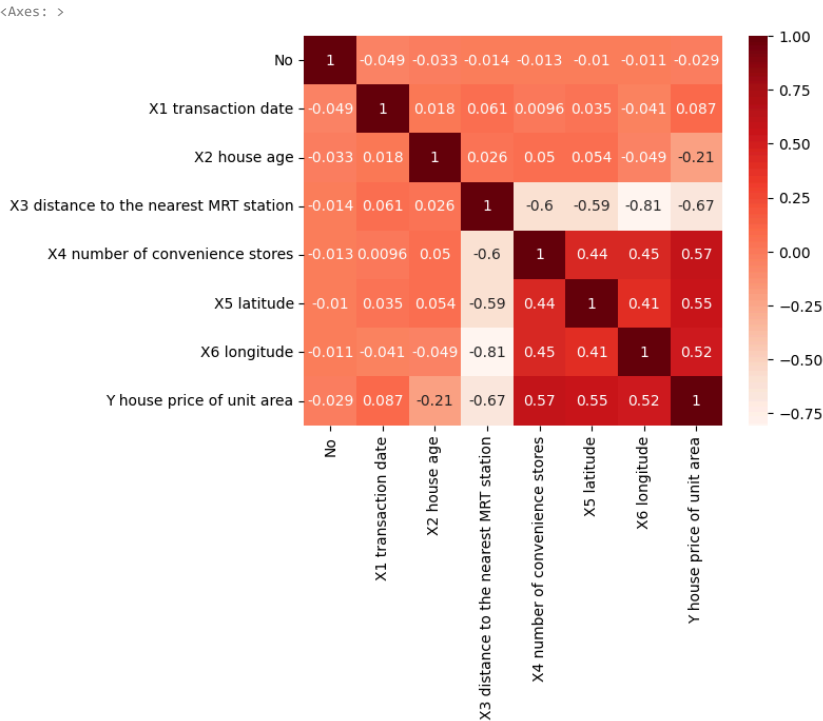
1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 414 entries, 0 to 413
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0    No                                    414 non-null    int64
1    X1 transaction date                  414 non-null    float64
2    X2 house age                        414 non-null    float64
3    X3 distance to the nearest MRT station 414 non-null    float64
4    X4 number of convenience stores       414 non-null    int64
5    X5 latitude                         414 non-null    float64
6    X6 longitude                        414 non-null    float64
7    Y house price of unit area           414 non-null    float64
dtypes: float64(6), int64(2)
memory usage: 26.0 KB
```

1 df.corr()

	No	X1 transaction date	X2 house age	X3 distance to the nearest MRT station	X4 number of convenience stores	X5 latitude	X6 longitude	Y house price of unit area
No	1.000000	-0.048658	-0.032808	-0.013573	-0.012699	-0.010110	-0.011059	-0.028587
X1 transaction date	-0.048658	1.000000	0.017549	0.060880	0.009635	0.035058	-0.041082	0.087491
X2 house age	-0.032808	0.017549	1.000000	0.025622	0.049593	0.054420	-0.048520	-0.210567
X3 distance to the nearest MRT station	-0.013573	0.060880	0.025622	1.000000	-0.602519	-0.591067	-0.806317	-0.673613
X4 number of convenience stores	-0.012699	0.009635	0.049593	-0.602519	1.000000	0.444143	0.449099	0.571005
X5 latitude	-0.010110	0.035058	0.054420	-0.591067	0.444143	1.000000	0.412924	0.546307

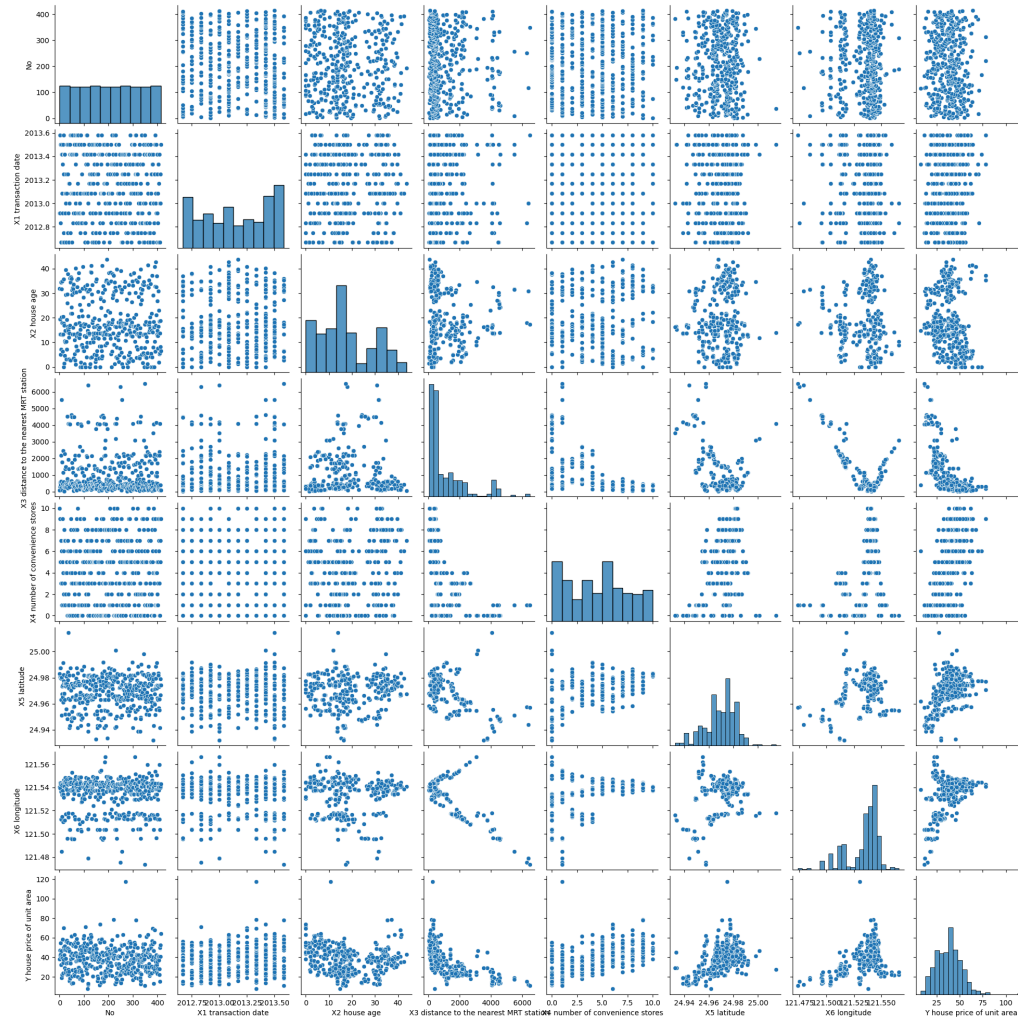
1 sns.heatmap(df.corr()),annot = True, cmap = 'Reds')



## Explanatory Data Analysis (EDA)

```
1 sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7a9c0fda8550>
```



## Training a Linear Regression Model

### X and Y arrays

```
1 X = df.drop('Y house price of unit area', axis = 1)
2 y = df['Y house price of unit area']
```

```
1 print("X = ",X.shape,"\ny = ",y.shape)
```

```
X = (414, 7)
y = (414,)
```

## Train Test Split

Now let's split the data into a training set and a testing set. We will train our model on the training set and then use the test set to evaluate the model.

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 101)
```

```
1 X_train.shape  
(289, 7)
```

```
1 X_test.shape  
(125, 7)
```

## Linear Regression

```
1 model = LinearRegression()
```

```
1 model.fit(X_train, y_train)
```

```
LinearRegression()
```

## Model Evaluation

```
1 model.coef_
```

```
array([-1.49344835e-17, -9.09342046e-15, -1.36338423e-16,  1.73472348e-18,  
       1.00000000e+00,  1.28927721e-14,  1.08238203e-14])
```

```
1 pd.DataFrame(model.coef_, X.columns, columns = ['Coefficients'])
```

	Coefficients
No	-1.493448e-17
X1 transaction date	-9.093420e-15
X2 house age	-1.363384e-16
X3 distance to the nearest MRT station	1.734723e-18
X4 number of convenience stores	1.000000e+00
X5 latitude	1.289277e-14
X6 longitude	1.082382e-14

## Predictions from our Model

```
1 y_pred = model.predict(X_test)
```

## Regression Evaluation Metrics

Here are three common evaluation metrics for regression problems:

- **Mean Absolute Error (MAE)** is the mean of the absolute value of the errors:

$$\frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

- **Mean Squared Error (MSE)** is the mean of the squared errors:

$$\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- **Root Mean Squared Error (RMSE)** is the square root of the mean of the mean of the squared errors:

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

🔴 Comparing these metrics:

- **MAE** is the easiest to understand, because it's the average error.
- **MSE** is more popular than MAE, because MSE "punishes" larger error, which tends to be useful in the real world.
- **RMSE** is even more popular than MSE, because RMSE is interpretable in the "y" units.

All of these are **loss functions**, because we want to minimize them.

```

1 MAE = metrics.mean_absolute_error(y_test, y_pred)
2 MSE = metrics.mean_squared_error(y_test, y_pred)
3 RMSE = np.sqrt(MSE)

```

```

1 MAE

4.231748536250847e-15

```

```

1 MSE

2.718688400256278e-29

```

```

1 RMSE

5.214104333685967e-15

```

```



1 df['X4 number of convenience stores'].mean()

4.094202898550725

```

## Residual Histogram

- Often for Linear Regression it is a good idea to separately evaluate residuals  $(y - \hat{y})$  and not just calculate the performance metrics(e.g. RMSE)
- Let's explore why this is important ...
- The residual errors should be random and close to a normal distribution



```

1 test_residual = y_test - y_pred

```

```


1 pd.DataFrame({'Error Values':(test_residual)}).plot.kde()

```

```


<Axes: ylabel='Density'>

```



```
1 sns.displot(test_residual, bins = 25 , kde =True)
```

<seaborn.axisgrid.FacetGrid at 0x7a9c0ba5a8c0>



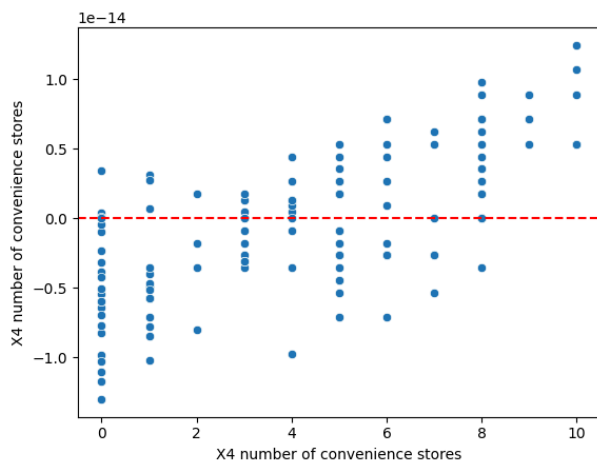
- Residual plot shows residual error VS. true y value

```
1 sns.scatterplot(x=y_test,y =test_residual)
```

```
2
```

```
3 plt.axhline(y=0,color = 'r', ls = '--')
```

<matplotlib.lines.Line2D at 0x7a9c08a67fa0>



- Residual plot showing clear pattern, indicating Linear Regression is no valid!