

---

**Name:** Jann Moises Nyll B. De los Reyes

**Section:** CPE22S3

**Date:** March 22, 2024

**Submitted to:** Engr. Roman M. Richard

---

## Logistic Regression Classifier Tutorial with Python

### 1. Introduction to Logistic Regression

When data scientist may come across a new classification problem, the first algorithm that may come across their mind is **Logistic Regression**. It is a supervised learning classification algorithm which is used to predict observations to a discrete set of classes. Practically, it is used to classify observations into different categories. Hence, its output is discrete in nature. **Logistic Regression** is also called **Logit Regression**. It is one of the most simple, straightforward and versatile classification algorithms which is used to solve classification problems.

### 2. Logistic Regression intuition

In statistics, the **Logistic Regression Model** is a widely used statistical model which is primarily used for classification purposes. It means that given a set of observations, Logistic Regression algorithm help us to classify these observations into two or more discrete classes. So, the target variable is discrete in nature.

The Logistic Regression algorithm works as follows -

#### Implement linear equation

Logistic Regression Algorithm works by implementing a linear equation with independent or explanatory variables to predict a response value. For example, we consider the example of number of hours studied and probability of passing the exam. Here, number of hour studied is the explanatory variable and it is denoted by  $x_1$ . Probability of passing the exam is the response or target variable and it is denoted by  $z$ .

If we have one explanatory variable( $x_1$ ) and one response variable( $z$ ), then linear equation would be given mathematically with the following equation-

$$z = \beta_0 + \beta_1 x_1$$

Here, the coefficient  $\beta_0$  and  $\beta_1$  are the parameters of the model.

If there are multiple explanatory variables, then the above equation can be extended to

$$z = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

Here, the coefficient  $\beta_0$ ,  $\beta_1$ ,  $\beta_2$  and  $\beta_n$  are the parameters of the model.

So, the predicted response value is given by the above equations and is denoted by  $z$ .

#### Sigmoid Function


This predicted response value, denoted by  $z$  is then converted into a probability value that lie between 0 and 1. We use the sigmoid function in order to map predicted values to probability values. This sigmoid function then maps any real value into a probability value between 0 and 1.

In machine learning, sigmoid function is used to map prediction to probabilities. The sigmoid function has an S shaped curve it is also called sigmoid curve.

A Sigmoid function is a special case of the Logistic function. It is given by the following mathematical formula.

Graphically, we can represent sigmoid function with the following graph.

#### Sigmoid Function



## Decision Boundary


The Sigmoid Function returns a probability between 0 and 1. This probability value is then mapped to a discrete class which is either "0" or "1". In order to map this probability value to a discrete class (pass/fail, yes/no, true/false), we select an threshold value. This threshold value is called Decision Boundary. Above this threshold value, we will map the probability values into class 1 and below which we will map values into class 0.

Mathematically, it can be expressed as follows:-

$$p \geq 0.5 \Rightarrow \text{class} = 1$$

$$p < 0.5 \Rightarrow \text{class} = 0$$

Generally, the decision boundary is set to 0.5, So, if the probability value is 0.8 ( $> 0.5$ ), we will map this observation to class 1. Similarly, if the probability value is 0.2 ( $< 0.5$ ), we will map this observation to class 0. This represent in the graph below -



## Making Predictions

Now, we know about sigmoid function and decision boundary in logistic regression. We can use our knowledge of sigmoid function and decision boundary to write a prediction function. A prediction function in logistic regression returns the probability of the observation being positive, YES or True. We call this as class 1 and it is denoted as  $P(\text{class} = 1)$ . If the probability inches closest to one, then we will be more confident about our model that the observation is in class 1, otherwise it is in class 0.

## 3 Assumptions of Logistic Regression

The Logistic Regression model requires several key assumptions. These are as follows:-

1. Logistic Regression model requires the dependent variable to be binary, multinomial or ordinal in nature.
2. It requires the observation to be independent to each other. So, the observation should not come from the repeated measurements.
3. Logistic Regression algorithm requires little or no multicollinearity among the independent variables. It means that the independent variables should not be too highly correlated with each other.

4. Logistic Regression algorithm requires model assumes linearity of independent variables and log odds.
5. The success of Logistic Regression model depends on the sample size. Typically requires a large sample size to achieve the high accuracy.

## 4 Types of Logistic Regression

Logistic Regression model can be classified into three groups based on the target variable categories. These three groups are describe below:-

### 1. Binary Logistic Regression

In Binary Logistic Regression, the target variable has two possible categories. The common examples of categories are yes or no, good or bad, true or false, spam or no spam and pass or fail.

### 2. Multinomial Logistic Regression

In Multinomial Logistic Regression, the target variable has three or more categories which are not in any particular order. So, there are three or more nominal categories. The examples include type of categories of fruits - apple, mango, orange and banana.

### 3. Ordinal Logistic Regression

In Ordinal Logistic Regression, the target variable has three or more ordinal categories. So, there is intrinsic order involved with the categories. For example, the student performance can be categories as poor, average, good and excellent.

## 5. Import libraries

```
In [111... #This python 3 environment comes with many helpful analytics libraries installed
#It is defined by the kaggle/python docker image: https://github.com/kaggle/docker-python
#For example, here's several helpful packages to Load in

import numpy as np #Linear algebra
import pandas as pd #data processing, csv file i/o (e.g. pd.read_csv)
import matplotlib.pyplot as plt #data visualization
import seaborn as sns #statistical data visualization
%matplotlib inline

#Input data files are available in the "../input/" directory
#For example, running this (by clicking run or pressing Shift + Enter) will list all files under the input directory

import os

for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

#Any result you write to the current directory are saved as output
```

```
In [112... import warnings
warnings.filterwarnings('ignore')
```

## 6.Import dataset

```
In [113... data = '/content/drive/MyDrive/Module 11/weatherAUS.csv'
df = pd.read_csv(data)
```

## 7. Explanatory Data Analysis

Now, we will explore the data to gain insights about the data

```
In [114... #view dimension of the dataset

df.shape
```

```
Out[114... (142193, 24)
```

We can see that there are 142193 instances and 24 variables in the data set.

```
In [115... # preview the dataset

df.head()
```

Out[115...

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity3pm
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	44.0	W	...	22.0
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	44.0	NNW	...	25.0
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	46.0	W	...	30.0
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	24.0	SE	...	16.0
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	41.0	ENE	...	33.0

5 rows × 24 columns



In [116...

```
col_names = df.columns  
  
col_names
```

Out[116...

```
Index(['Date', 'Location', 'MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation',  
      'Sunshine', 'WindGustDir', 'WindGustSpeed', 'WindDir9am', 'WindDir3pm',  
      'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm',  
      'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am',  
      'Temp3pm', 'RainToday', 'RISK_MM', 'RainTomorrow'],  
      dtype='object')
```

Drop RISK\_MM variable

It is given in the dataset description, that we should drop the RISK\_MM feature variable from the dataset description. So, we should drop it as follows-

In [117...

```
df.drop(['RISK_MM'],axis =1, inplace = True)
```

In [118...

```
#view summary of dataset  
  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 142193 entries, 0 to 142192  
Data columns (total 23 columns):  
#   Column          Non-Null Count  Dtype  
---  ---  
0   Date            142193 non-null object  
1   Location        142193 non-null object  
2   MinTemp         141556 non-null float64  
3   MaxTemp         141871 non-null float64  
4   Rainfall        140787 non-null float64  
5   Evaporation     81350 non-null float64  
6   Sunshine        74377 non-null float64  
7   WindGustDir     132863 non-null object  
8   WindGustSpeed   132923 non-null float64  
9   WindDir9am      132180 non-null object  
10  WindDir3pm      138415 non-null object  
11  WindSpeed9am    140845 non-null float64  
12  WindSpeed3pm    139563 non-null float64  
13  Humidity9am     140419 non-null float64  
14  Humidity3pm     138583 non-null float64  
15  Pressure9am     128179 non-null float64  
16  Pressure3pm     128212 non-null float64  
17  Cloud9am        88536 non-null float64  
18  Cloud3pm        85099 non-null float64  
19  Temp9am         141289 non-null float64  
20  Temp3pm         139467 non-null float64  
21  RainToday       140787 non-null object  
22  RainTomorrow    142193 non-null object  
dtypes: float64(16), object(7)  
memory usage: 25.0+ MB
```

Types of variables

In this section, I segregate the dataset into categorical and numerical variables. There are a mixture of categorical and numerical variables in the dataset. Categorical variables have data type object. Numerical variables have data type float64.

first of all, we will find categorical variables

```
In [119... # find categorical variables

categorical = [var for var in df.columns if df[var].dtype == 'O']

print('There are {} categorical variables\n'.format(len(categorical)))
print('The categorical variables are:', categorical)
```

There are 7 categorical variables

The categorical variables are: ['Date', 'Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday', 'RainTomorrow']

```
In [120... #view categorical variables

df[categorical].head()
```

```
Out[120...      Date  Location  WindGustDir  WindDir9am  WindDir3pm  RainToday  RainTomorrow
0  2008-12-01   Albury           W           W          WNW         No           No
1  2008-12-02   Albury        WNW          NNW          WSW         No           No
2  2008-12-03   Albury        WSW           W          WSW         No           No
3  2008-12-04   Albury          NE           SE           E         No           No
4  2008-12-05   Albury           W          ENE          NW         No           No
```

## Summary of categorical variables

- There is a date variable. It is denoted by `Date` column.
- There are 6 categorical variables. These are given by `Location`, `WindGustDir`, `WindDir9am`, `WindDir3pm`, `RainToday` and `RainTomorrow`.
- There are two binary categorical variables - `RainToday` and `RainTomorrow`.
- `RainTomorrow` is the target variable.

## Explore problems within categorical variables

First, I will explore the categorical variables.

## Missing Values in categorical variables

```
In [121... #Check missing values in categorical variables

df[categorical].isnull().sum()
```

```
Out[121... Date           0
Location          0
WindGustDir       9330
WindDir9am       10013
WindDir3pm        3778
RainToday        1406
RainTomorrow      0
dtype: int64
```

```
In [122... #print categorical variables containing missing values

cat1 = [var for var in categorical if df[var].isnull().sum() != 0]

print(df[cat1].isnull().sum())
```

```
WindGustDir    9330
WindDir9am     10013
WindDir3pm      3778
RainToday      1406
dtype: int64
```

We can see that there are only 4 categorical variables in dataset which contains missing values. These are `WindGustDir`, `WindDir9am`, `WindDir3pm` and `RainToday`.

## Frequency counts of categorical variables

Now, I will check the frequency counts of categorical variables.

```
In [123... # view frequency of categorical variables

for var in categorical:

    print(df[var].value_counts()/float(len(df)))
```

```

Date
2013-12-01    0.000345
2014-01-09    0.000345
2014-01-11    0.000345
2014-01-12    0.000345
2014-01-13    0.000345
...
2007-11-29    0.000007
2007-11-28    0.000007
2007-11-27    0.000007
2007-11-26    0.000007
2008-01-31    0.000007

```

Name: count, Length: 3436, dtype: float64

```

Location
Canberra      0.024038
Sydney        0.023468
Perth         0.022455
Darwin        0.022448
Hobart        0.022420
Brisbane      0.022230
Adelaide      0.021731
Bendigo       0.021337
Townsville    0.021330
AliceSprings  0.021316
MountGambier  0.021309
Launceston    0.021295
Ballarat      0.021295
Albany        0.021211
Albury        0.021175
PerthAirport  0.021161
MelbourneAirport 0.021161
Mildura       0.021147
SydneyAirport 0.021133
Nuriootpa     0.021112
Sale          0.021098
Watsonia      0.021091
Tuggeranong   0.021084
Portland      0.021070
Woomera       0.021028
Cairns        0.021014
Cobar         0.021014
Wollongong    0.020979
GoldCoast     0.020957
WaggaWagga    0.020929
Penrith       0.020845
NorfolkIsland 0.020845
SalmonGums    0.020782
Newcastle     0.020782
CoffsHarbour  0.020768
Witchcliffe   0.020761
Richmond      0.020753
Dartmoor      0.020697
NorahHead     0.020599
BadgerysCreek 0.020592
MountGinini   0.020444
Moree         0.020071
Walpole       0.019825
PearceRAAF    0.019424
Williamstown  0.017954
Melbourne     0.017125
Nhil          0.011034
Katherine     0.010964
Uluru         0.010697

```

Name: count, dtype: float64

```

WindGustDir
W      0.068780
SE     0.065467
E      0.063794
N      0.063526
SSE    0.063245
S       0.062936
WSW    0.062598
SW      0.061867
SSW    0.060552
WNW    0.056726
NW      0.056283
ENE     0.056205
ESE     0.051374
NE      0.049651
NNW     0.046142

```

```

NNE      0.045241
Name: count, dtype: float64
WindDir9am
N         0.080123
SE        0.064434
E         0.063463
SSE       0.063055
NW        0.060144
S         0.059729
W         0.058090
SW        0.057928
NNE       0.055896
NNW       0.055136
ENE       0.054398
ESE       0.053153
NE        0.052935
SSW       0.052380
WNW       0.050593
WSW       0.048125
Name: count, dtype: float64
WindDir3pm
SE        0.074990
W         0.069701
S         0.067500
WSW       0.065608
SW        0.064574
SSE       0.064293
N         0.060952
WNW       0.060875
NW        0.059553
ESE       0.058948
E         0.058667
NE        0.057415
SSW       0.056332
NNW       0.054384
ENE       0.054321
NNE       0.045319
Name: count, dtype: float64
RainToday
No        0.768899
Yes       0.221213
Name: count, dtype: float64
RainTomorrow
No        0.775819
Yes       0.224181
Name: count, dtype: float64

```

## Number of labels: cardinality

The number of labels within a categorical variables is known as **cardinality**. A high number of labels within a variable is known as **high cardinality**. High cardinality may pose some serious problems in the machine learning model. So, we will check for high cardinality.

```

In [124...] #check for cardinality in categorical variables

for var in categorical:

    print(var, ' contains ', len(df[var].unique()), ' labels')

```

```

Date contains 3436 labels
Location contains 49 labels
WindGustDir contains 17 labels
WindDir9am contains 17 labels
WindDir3pm contains 17 labels
RainToday contains 3 labels
RainTomorrow contains 2 labels

```

We can see that there is a `Date` variable which needs to be preprocessed. We will do preprocessing in the following section.

All the other variables contain relatively smaller number of variables.

## Feature Engineering of Date Variable

```

In [125...] df['Date'].dtypes

Out[125...] dtype('O')

```

We can see that the data type of `Date` variable is object. We will parse the date currently coded as object into datetime format.



```
In [126... # parse the dates, currently coded as string, into datetime format
```

```
df['Date'] = pd.to_datetime(df['Date'])
```

```
In [127... #extract year from date
```

```
df['Year'] = df['Date'].dt.year
```

```
df['Year'].head()
```

```
Out[127... 0    2008
1    2008
2    2008
3    2008
4    2008
Name: Year, dtype: int32
```

```
In [128... #extract year from month
```

```
df['Month'] = df['Date'].dt.month
```

```
df['Month'].head()
```

```
Out[128... 0    12
1    12
2    12
3    12
4    12
Name: Month, dtype: int32
```

```
In [129... #extract year from day
```

```
df['Day'] = df['Date'].dt.day
```

```
df['Day'].head()
```

```
Out[129... 0    1
1    2
2    3
3    4
4    5
Name: Day, dtype: int32
```

```
In [130... #again view the summary of Dataset
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 142193 entries, 0 to 142192
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  142193 non-null  datetime64[ns]
1   Location              142193 non-null  object
2   MinTemp               141556 non-null  float64
3   MaxTemp               141871 non-null  float64
4   Rainfall              140787 non-null  float64
5   Evaporation           81350 non-null   float64
6   Sunshine              74377 non-null   float64
7   WindGustDir           132863 non-null  object
8   WindGustSpeed         132923 non-null  float64
9   WindDir9am            132180 non-null  object
10  WindDir3pm            138415 non-null  object
11  WindSpeed9am          140845 non-null  float64
12  WindSpeed3pm          139563 non-null  float64
13  Humidity9am           140419 non-null  float64
14  Humidity3pm           138583 non-null  float64
15  Pressure9am           128179 non-null  float64
16  Pressure3pm           128212 non-null  float64
17  Cloud9am              88536 non-null   float64
18  Cloud3pm              85099 non-null   float64
19  Temp9am               141289 non-null  float64
20  Temp3pm               139467 non-null  float64
21  RainToday             140787 non-null  object
22  RainTomorrow          142193 non-null  object
23  Year                  142193 non-null  int32
24  Month                 142193 non-null  int32
25  Day                   142193 non-null  int32
dtypes: datetime64[ns](1), float64(16), int32(3), object(6)
memory usage: 26.6+ MB
```

We can see that there are three additional columns created from `Date` variable. Now, I will drop the original `Date` variable from the dataset.

```
In [131... #drop the original Date variable
df.drop('Date',axis = 1, inplace= True)
```

```
In [132... #preview the dataset again
df.head()
```

```
Out[132...      Location  MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  WindGustDir  WindGustSpeed  WindDir9am  WindDir3pm  ...  Pressure
0    Albury      13.4     22.9      0.6        NaN        NaN           W           44.0           W        WNW      ...        1
1    Albury       7.4     25.1      0.0        NaN        NaN        WNW           44.0          NNW        WSW      ...        1
2    Albury      12.9     25.7      0.0        NaN        NaN        WSW           46.0           W        WSW      ...        1
3    Albury       9.2     28.0      0.0        NaN        NaN          NE           24.0           SE           E      ...        1
4    Albury      17.5     32.3      1.0        NaN        NaN           W           41.0          ENE         NW      ...        1
```

5 rows × 25 columns



Now, we can see that the `Date` variable has been removed to the dataset.

## Explore Categorical Variables

Now, we will explore the categorical variables one by one.

```
In [133... #find categorical variables

categorical = [var for var in df.columns if df[var].dtypes == 'O']

print('There are {} categorical variables\n'.format(len(categorical)))

print('The categorical variables are:', categorical)
```

There are 6 categorical variables

The categorical variables are: ['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday', 'RainTomorrow']

We can see that there are 6 categorical variables in the dataset. The `Date` variable has been removed. First, I will check missing values in categorical variables.

```
In [134... #check for missing values in categorical variables

df[categorical].isnull().sum()
```

```
Out[134... Location      0
WindGustDir    9330
WindDir9am     10013
WindDir3pm     3778
RainToday      1406
RainTomorrow    0
dtype: int64
```

We can see that `WindGustDir`, `WindDir9am`, `WindDir3pm`, `RainToday` variables contain missing values. We will explore these variables one by one.

## Explore Location variable

```
In [135... #print number of labels in Location variable

print('Location contains', len(df.Location.unique()), 'labels')
```

Location contains 49 labels

```
In [136... #check labels in Location variable

df.Location.unique()
```

```
Out[136...] array(['Albury', 'BadgerysCreek', 'Cobar', 'CoffsHarbour', 'Moree',
      'Newcastle', 'NorahHead', 'NorfolkIsland', 'Penrith', 'Richmond',
      'Sydney', 'SydneyAirport', 'WaggaWagga', 'Williamstown',
      'Wollongong', 'Canberra', 'Tuggeranong', 'MountGinini', 'Ballarat',
      'Bendigo', 'Sale', 'MelbourneAirport', 'Melbourne', 'Mildura',
      'Nhil', 'Portland', 'Watsonia', 'Dartmoor', 'Brisbane', 'Cairns',
      'GoldCoast', 'Townsville', 'Adelaide', 'MountGambier', 'Nuriootpa',
      'Woomera', 'Albany', 'Witchcliffe', 'PearceRAAF', 'PerthAirport',
      'Perth', 'SalmonGums', 'Walpole', 'Hobart', 'Launceston',
      'AliceSprings', 'Darwin', 'Katherine', 'Uluru'], dtype=object)
```

```
In [137...] #check the frequency distribution of values in Location variable
```

```
df.Location.value_counts()
```

```
Out[137...] Location
Canberra      3418
Sydney        3337
Perth         3193
Darwin        3192
Hobart        3188
Brisbane      3161
Adelaide      3090
Bendigo       3034
Townsville    3033
AliceSprings  3031
MountGambier  3030
Launceston    3028
Ballarat      3028
Albany        3016
Albury        3011
PerthAirport  3009
MelbourneAirport 3009
Mildura       3007
SydneyAirport 3005
Nuriootpa     3002
Sale          3000
Watsonia      2999
Tuggeranong   2998
Portland      2996
Woomera       2990
Cairns        2988
Cobar         2988
Wollongong    2983
GoldCoast     2980
WaggaWagga    2976
Penrith       2964
NorfolkIsland 2964
SalmonGums    2955
Newcastle     2955
CoffsHarbour  2953
Witchcliffe   2952
Richmond     2951
Dartmoor      2943
NorahHead     2929
BadgerysCreek 2928
MountGinini   2907
Moree         2854
Walpole       2819
PearceRAAF    2762
Williamstown  2553
Melbourne     2435
Nhil          1569
Katherine     1559
Uluru         1521
Name: count, dtype: int64
```

```
In [138...] #Let's Do One Hot Encoding of Location variable
#get k-1 dummy variables after One Hot Encoding
#preview the dataset with head() method
```

```
pd.get_dummies(df.Location, drop_first = True).astype(int).head()
```

Out[138...

	Albany	Albury	AliceSprings	BadgerysCreek	Ballarat	Bendigo	Brisbane	Cairns	Canberra	Cobar	...	Townsville	Tuggeranong	Ulur
0	0	1	0	0	0	0	0	0	0	0	...	0	0	
1	0	1	0	0	0	0	0	0	0	0	...	0	0	
2	0	1	0	0	0	0	0	0	0	0	...	0	0	
3	0	1	0	0	0	0	0	0	0	0	...	0	0	
4	0	1	0	0	0	0	0	0	0	0	...	0	0	

5 rows × 48 columns



Explore WindGustDir variable

In [139...

```
#print number of Labels in WindGustDir variable
print('WindGustDir contains', len(df.WindGustDir.unique()), 'labels')
```

WindGustDir contains 17 labels

In [140...

```
#check Labels in WindGustDir variable
df['WindGustDir'].unique()
```

Out[140...

```
array(['W', 'WNW', 'WSW', 'NE', 'NNW', 'N', 'NNE', 'SW', 'ENE', 'SSE',
       'S', 'NW', 'SE', 'ESE', nan, 'E', 'SSW'], dtype=object)
```

In [141...

```
#check frequency distribution of values in WindGustDir variable
df.WindGustDir.value_counts()
```

Out[141...

```
WindGustDir
W      9780
SE     9309
E      9071
N      9033
SSE    8993
S      8949
WSW    8901
SW     8797
SSW    8610
WNW    8066
NW     8003
ENE    7992
ESE    7305
NE     7060
NNW    6561
NNE    6433
Name: count, dtype: int64
```

In [142...

```
#Let's Do One Hot Encoding of WindGustDir variable
#get k-1 dummy variables after One Hot Encoding
#also add an additional dummy variable to indicate there was missing data
#preview the dataset with head() method

pd.get_dummies(df.Location, drop_first = True, dummy_na = True).astype(int).head()
```

Out[142...

	Albany	Albury	AliceSprings	BadgerysCreek	Ballarat	Bendigo	Brisbane	Cairns	Canberra	Cobar	...	Tuggeranong	Uluru	WaggaW
0	0	1	0	0	0	0	0	0	0	0	...	0	0	
1	0	1	0	0	0	0	0	0	0	0	...	0	0	
2	0	1	0	0	0	0	0	0	0	0	...	0	0	
3	0	1	0	0	0	0	0	0	0	0	...	0	0	
4	0	1	0	0	0	0	0	0	0	0	...	0	0	

5 rows × 49 columns



In [144...

```
#sum the number for 1s per boolean variable over the rows of the dataset
#it will tell us how many observation we have for each category
```

```
pd.get_dummies(df.WindGustDir, drop_first = True, dummy_na = True).astype(int).sum(axis = 0)
```

```
Out[144...  ENE    7992
           ESE    7305
           N     9033
           NE     7060
           NNE    6433
           NNW    6561
           NW     8003
           S     8949
           SE     9309
           SSE    8993
           SSW    8610
           SW     8797
           W     9780
           WNW    8066
           WSW    8901
           NaN    9330
dtype: int64
```

We can see that there are 9330 missing values in WindGustDir variable.

## Explore WindDir9am variable

```
In [145... # print number of labels in WindDir9am variable

print('WindDir9am contains', len(df['WindDir9am'].unique()), 'labels')
```

WindDir9am contains 17 labels

```
In [146... #check labels in WindDir9am variable

df['WindDir9am'].unique()
```

```
Out[146... array(['W', 'NNW', 'SE', 'ENE', 'SW', 'SSE', 'S', 'NE', nan, 'SSW', 'N',
        'WSW', 'ESE', 'E', 'NW', 'WNW', 'NNE'], dtype=object)
```

```
In [147... #check frequency distribution of values in WindDir9am variable

df['WindDir9am'].value_counts()
```

```
Out[147... WindDir9am
N          11393
SE          9162
E           9024
SSE         8966
NW          8552
S           8493
W           8260
SW          8237
NNE         7948
NNW         7840
ENE         7735
ESE         7558
NE          7527
SSW         7448
WNW         7194
WSW         6843
Name: count, dtype: int64
```

```
In [148... #Let's Do One Hot Encoding of WindDir9am variable
#get k-1 dummy variables after One Hot Encoding
#also add an additional dummy variable to indicate there was missing data
#preview the dataset with head() method

pd.get_dummies(df.WindDir9am, drop_first = True, dummy_na = True).astype(int).head()
```

```
Out[148...
```

	ENE	ESE	N	NE	NNE	NNW	NW	S	SE	SSE	SSW	SW	W	WNW	WSW	NaN
0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0
1	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0
3	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0
4	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

```
In [149... #sum the number for 1s per boolean variable over the rows of the dataset
#it will tell us how many observation we have for each category

pd.get_dummies(df.WindDir9am, drop_first = True, dummy_na = True).astype(int).sum(axis = 0)
```

```
Out[149... ENE      7735
ESE      7558
N       11393
NE       7527
NNE      7948
NNW      7840
NW       8552
S        8493
SE       9162
SSE      8966
SSW      7448
SW       8237
W        8260
WNW      7194
WSW      6843
NaN     10013
dtype: int64
```

## Explore WindDir3pm variable

```
In [150... # print number of labels in WindDir3pm variable

print('WindDir3pm contains', len(df['WindDir3pm'].unique()), 'labels')
```

WindDir3pm contains 17 labels

```
In [151... #check labels in WindDir3pm variable

df['WindDir3pm'].unique()
```

```
Out[151... array(['WNW', 'WSW', 'E', 'NW', 'W', 'SSE', 'ESE', 'ENE', 'NNW', 'SSW',
      'SW', 'SE', 'N', 'S', 'NNE', nan, 'NE'], dtype=object)
```

```
In [152... #check frequency distribution of values in WindDir3pm variable

df['WindDir3pm'].value_counts()
```

```
Out[152... WindDir3pm
SE      10663
W       9911
S       9598
WSW     9329
SW      9182
SSE     9142
N       8667
WNW     8656
NW      8468
ESE     8382
E       8342
NE      8164
SSW     8010
NNW     7733
ENE     7724
NNE     6444
Name: count, dtype: int64
```

```
In [153... #Let's Do One Hot Encoding of WindDir3pm variable
#get k-1 dummy variables after One Hot Encoding
#also add an additional dummy variable to indicate there was missing data
#preview the dataset with head() method
```

```
pd.get_dummies(df.WindDir3pm, drop_first = True, dummy_na = True).astype(int).head()
```

Out[153...

	ENE	ESE	N	NE	NNE	NNW	NW	S	SE	SSE	SSW	SW	W	WNW	WSW	NaN
0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0
2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0
3	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0

In [154...

```
#sum the number for 1s per boolean variable over the rows of the dataset
#it will tell us how many observation we have for each category
```

```
pd.get_dummies(df.WindDir3pm, drop_first = True, dummy_na = True).astype(int).sum(axis = 0)
```

Out[154...

```
ENE      7724
ESE      8382
N        8667
NE       8164
NNE      6444
NNW      7733
NW       8468
S        9598
SE      10663
SSE      9142
SSW      8010
SW       9182
W        9911
WNW      8656
WSW      9329
NaN      3778
dtype: int64
```

There are 3778 missing values in the `WindDir3pm` variable.

## Explore RainToday variable

In [155...

```
# print number of labels in RainToday variable

print('RainToday contains', len(df['RainToday'].unique()), 'labels')
```

RainToday contains 3 labels

In [156...

```
#check labels in RainToday variable

df['RainToday'].unique()
```

Out[156...

```
array(['No', 'Yes', nan], dtype=object)
```

In [157...

```
#check frequency distribution of values in RainToday variable

df['RainToday'].value_counts()
```

Out[157...

```
RainToday
No      109332
Yes      31455
Name: count, dtype: int64
```

In [158...

```
#Let's Do One Hot Encoding of RainToday variable
#get k-1 dummy variables after One Hot Encoding
#also add an additional dummy variable to indicate there was missing data
#preview the dataset with head() method

pd.get_dummies(df.RainToday, drop_first = True, dummy_na = True).astype(int).head()
```

```
Out[158...
   Yes  NaN
0    0    0
1    0    0
2    0    0
3    0    0
4    0    0
```

```
In [159...
#sum the number for 1s per boolean variable over the rows of the dataset
#it will tell us how many observation we have for each category

pd.get_dummies(df.RainToday, drop_first = True, dummy_na = True).astype(int).sum(axis = 0)
```

```
Out[159...
Yes      31455
NaN       1406
dtype: int64
```

There are 1406 missing values in the `RainToday` variable.

## Explore Numerical Variables

```
In [160...
#find numerical variables

numerical = [var for var in df.columns if df[var].dtypes != 'O']

print('There are {} numerical variables\n'.format(len(numerical)))

print('The numerical variables are: ',numerical)
```

There are 19 numerical variables

The numerical variables are: ['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine', 'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm', 'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am', 'Temp3pm', 'Year', 'Month', 'Day']

```
In [161...
#view the numerical values

df[numerical].head()
```

```
Out[161...
   MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  WindGustSpeed  WindSpeed9am  WindSpeed3pm  Humidity9am  Humidity3pm
0      13.4      22.9        0.6          NaN        NaN             44.0             20.0             24.0             71.0             22.0
1       7.4      25.1         0.0          NaN        NaN             44.0              4.0             22.0             44.0             25.0
2      12.9      25.7         0.0          NaN        NaN             46.0             19.0             26.0             38.0             30.0
3       9.2      28.0         0.0          NaN        NaN             24.0             11.0              9.0             45.0             16.0
4      17.5      32.3         1.0          NaN        NaN             41.0              7.0             20.0             82.0             33.0
```



### Summary of numerical variables

- There are 16 numerical variables.
- These are given by `MinTemp`, `MaxTemp`, `Rainfall`, `Evaporation`, `Sunshine`, `WindGustSpeed`, `WindSpeed9am`, `WindSpeed3pm`, `Humidity9am`, `Humidity3pm`, `Pressure9am`, `Pressure3pm`, `Cloud9am`, `Cloud3pm`, `Temp9am`, and `Temp3pm`.
- All of the numerical variable are of continuous type.

## Explore problems within numerical variables

Now, We will explore the numerical variables

### Missing values in numerical variables

```
In [162...
#check missing values in numerical variables

df[numerical].isnull().sum()
```



```
Out[162...  MinTemp      637
             MaxTemp      322
             Rainfall     1406
             Evaporation   60843
             Sunshine     67816
             WindGustSpeed  9270
             WindSpeed9am  1348
             WindSpeed3pm  2630
             Humidity9am   1774
             Humidity3pm   3610
             Pressure9am   14014
             Pressure3pm   13981
             Cloud9am     53657
             Cloud3pm     57094
             Temp9am      904
             Temp3pm     2726
             Year         0
             Month        0
             Day          0
             dtype: int64
```

We can see that all the 16 numerical variables contain missing values.

## Outliers in numerical variables

```
In [163... # view summary statistics in numerical variables

print(round(df[numerical].describe()),2)
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	\
count	141556.0	141871.0	140787.0	81350.0	74377.0	132923.0	
mean	12.0	23.0	2.0	5.0	8.0	40.0	
std	6.0	7.0	8.0	4.0	4.0	14.0	
min	-8.0	-5.0	0.0	0.0	0.0	6.0	
25%	8.0	18.0	0.0	3.0	5.0	31.0	
50%	12.0	23.0	0.0	5.0	8.0	39.0	
75%	17.0	28.0	1.0	7.0	11.0	48.0	
max	34.0	48.0	371.0	145.0	14.0	135.0	

	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	Pressure9am	\
count	140845.0	139563.0	140419.0	138583.0	128179.0	
mean	14.0	19.0	69.0	51.0	1018.0	
std	9.0	9.0	19.0	21.0	7.0	
min	0.0	0.0	0.0	0.0	980.0	
25%	7.0	13.0	57.0	37.0	1013.0	
50%	13.0	19.0	70.0	52.0	1018.0	
75%	19.0	24.0	83.0	66.0	1022.0	
max	130.0	87.0	100.0	100.0	1041.0	

	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	Year	\
count	128212.0	88536.0	85099.0	141289.0	139467.0	142193.0	
mean	1015.0	4.0	5.0	17.0	22.0	2013.0	
std	7.0	3.0	3.0	6.0	7.0	3.0	
min	977.0	0.0	0.0	-7.0	-5.0	2007.0	
25%	1010.0	1.0	2.0	12.0	17.0	2011.0	
50%	1015.0	5.0	5.0	17.0	21.0	2013.0	
75%	1020.0	7.0	7.0	22.0	26.0	2015.0	
max	1040.0	9.0	9.0	40.0	47.0	2017.0	

	Month	Day
count	142193.0	142193.0
mean	6.0	16.0
std	3.0	9.0
min	1.0	1.0
25%	3.0	8.0
50%	6.0	16.0
75%	9.0	23.0
max	12.0	31.0

On closer inspection, we can see that the `Rainfall`, `Evaporation`, `WindSpeed9am` and `WindSpeed3pm` columns may contain outliers.

We will draw boxplot to visualize outliers in above variables

```
In [164... # draw boxplot to visualize outliers

plt.figure(figsize=(15,10))

plt.subplot(2,2,1)
fig = df.boxplot(column = 'Rainfall')
```

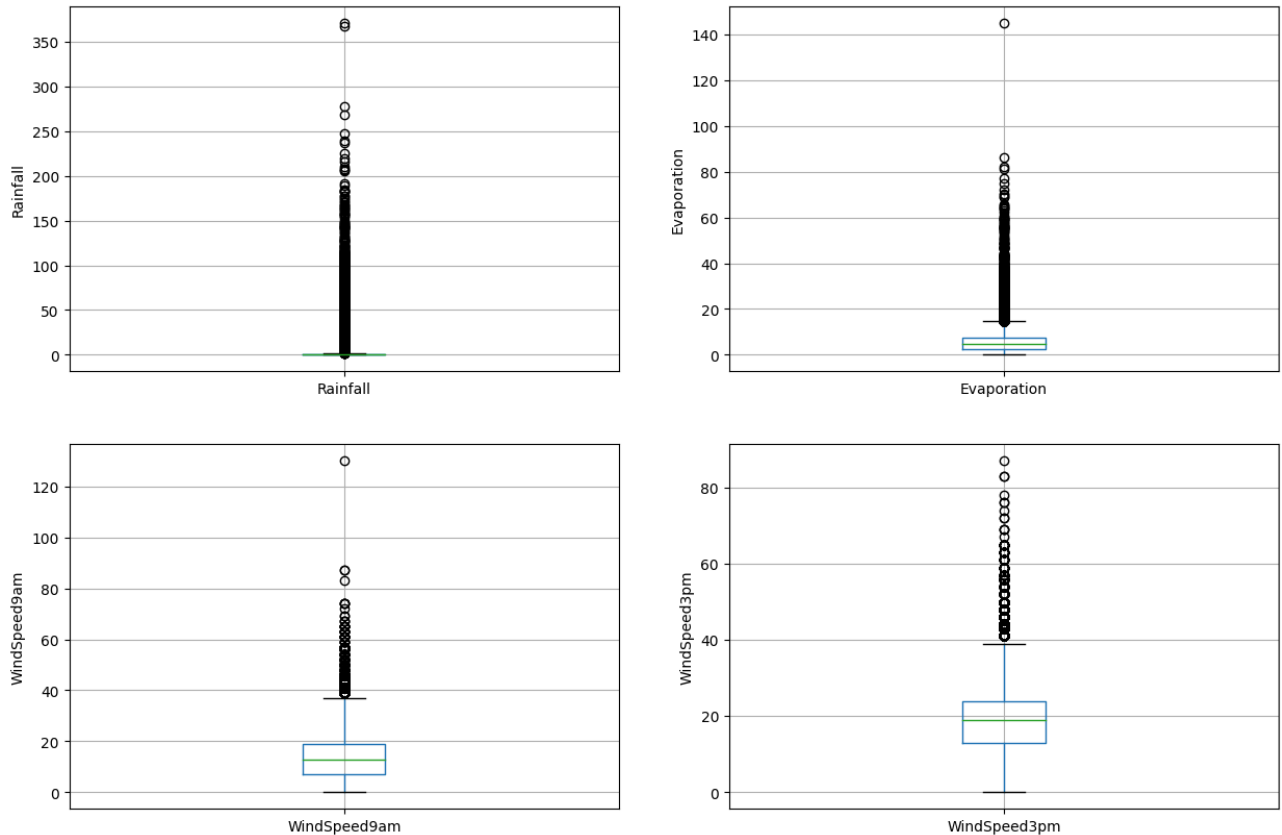
```
fig.set_title('')
fig.set_ylabel('Rainfall')

plt.subplot(2,2,2)
fig = df.boxplot(column = 'Evaporation')
fig.set_title('')
fig.set_ylabel('Evaporation')

plt.subplot(2,2,3)
fig = df.boxplot(column = 'WindSpeed9am')
fig.set_title('')
fig.set_ylabel('WindSpeed9am')

plt.subplot(2,2,4)
fig = df.boxplot(column = 'WindSpeed3pm')
fig.set_title('')
fig.set_ylabel('WindSpeed3pm')
```

Out[164...] Text(0, 0.5, 'WindSpeed3pm')



The above boxplot confirm that there are lot of outliers in these variables.

## Check the distribution of variables

Now, I will plot the histogram to check distribution to find out if they are normal or skewed. If the variable follows normal distribution, then we will do Extreme Value Analysis otherwise if they are skewed, We will find IQR (Interquartile range)

In [165...] #Plot histogram to check distribution

```
plt.figure(figsize=(15,10))

plt.subplot(2,2,1)
fig = df.Rainfall.hist(bins=10)
fig.set_xlabel('Rainfall')
fig.set_ylabel('RainTommorow')





plt.subplot(2,2,2)
fig = df.Evaporation.hist(bins=10)
fig.set_xlabel('Evaporation')
fig.set_ylabel('RainTommorow')

plt.subplot(2,2,3)
fig = df.WindSpeed9am.hist(bins=10)
```

```
fig.set_xlabel('WindSpeed9am')
fig.set_ylabel('RainTomorrow')

plt.subplot(2,2,4)
fig = df.WindSpeed3pm.hist(bins=10)
fig.set_xlabel('WindSpeed3pm')
fig.set_ylabel('RainTomorrow')
```

Out[165... Text(0, 0.5, 'RainTomorrow')



We can see that all the four variables are skewed. So, we will use interquartile range to find outliers.

```
In [166... #find outliers for Rainfall variable

IQR = df.Rainfall.quantile(0.75)- df.Rainfall.quantile(0.25)
Lower_fence = df.Rainfall.quantile(0.25)-(IQR*3)
Upper_fence = df.Rainfall.quantile(0.75)+(IQR*3)
```

```
print('Rainfall outliers are values < {lowerboundary} or > {upperboundary}').format(lowerboundary=Lower_fence, upperboundary =U
```

Rainfall outliers are values < -2.4000000000000004 or > 3.2

For `Rainfall`, the minimum and maximum values are 0.0 and 371.0. So, the outliers are values > 3.2.

```
In [167... #find outliers for Evaporation variable

IQR = df.Evaporation.quantile(0.75)- df.Evaporation.quantile(0.25)
Lower_fence = df.Evaporation.quantile(0.25)-(IQR*3)
Upper_fence = df.Evaporation.quantile(0.75)+(IQR*3)
```

```
print('Evaporation outliers are values < {lowerboundary} or > {upperboundary}').format(lowerboundary=Lower_fence, upperboundary
```

Evaporation outliers are values < -11.800000000000002 or > 21.800000000000004

For `Evaporation`, the minimum and maximum values are 0.0 and 145.0. So, the outliers are values > 21.8.

```
In [168... #find outliers for WindSpeed9am variable

IQR = df.WindSpeed9am.quantile(0.75)- df.WindSpeed9am.quantile(0.25)
Lower_fence = df.WindSpeed9am.quantile(0.25)-(IQR*3)
Upper_fence = df.WindSpeed9am.quantile(0.75)+(IQR*3)
```

```
print('WindSpeed9am outliers are values < {lowerboundary} or > {upperboundary}').format(lowerboundary=Lower_fence, upperboundar
```

WindSpeed9am outliers are values < -29.0 or > 55.0

For `WindSpeed9am`, the minimum and maximum values are 0.0 and 130.0. So, the outliers are values > 55.0.

```
In [169... #find outliers for WindSpeed3pm variable

IQR = df.WindSpeed3pm.quantile(0.75)- df.WindSpeed3pm.quantile(0.25)
Lower_fence = df.WindSpeed3pm.quantile(0.25)-(IQR*3)
Upper_fence = df.WindSpeed3pm.quantile(0.75)+(IQR*3)

print('WindSpeed3pm outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerboundary=Lower_fence, upperboundar
```

WindSpeed3pm outliers are values < -20.0 or > 57.0

For `WindSpeed3pm`, the minimum and maximum values are 0.0 and 87.0. So, the outliers are values > 57.0.

## 8. Declare feature vector and target variable

```
In [170... X = df.drop(['RainTomorrow'],axis = 1)
y = df['RainTomorrow']
```

## 9. Split data into separate training and test set

```
In [171... #split X and y into training and testing sets

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
```

```
In [172... #check the shape of X_train and X_test

X_train.shape, X_test.shape
```

```
Out[172... ((113754, 24), (28439, 24))
```

## 10. Feature Engineering

**Feature Engineering** is the process of transforming raw data into useful features that help us to understand our model better and increase its predictive power. We will carry out feature engineering on different types of variables.

First, we will display the categorical and numerical variables again separately.

```
In [173... #check data types in X_train

X_train.dtypes
```

```
Out[173... Location          object
MinTemp           float64
MaxTemp           float64
Rainfall          float64
Evaporation       float64
Sunshine          float64
WindGustDir       object
WindGustSpeed     float64
WindDir9am        object
WindDir3pm        object
WindSpeed9am      float64
WindSpeed3pm      float64
Humidity9am       float64
Humidity3pm       float64
Pressure9am       float64
Pressure3pm       float64
Cloud9am          float64
Cloud3pm          float64
Temp9am           float64
Temp3pm           float64
RainToday         object
Year              int32
Month             int32
Day               int32
dtype: object
```

```
In [174... #display categorical variables

categorical = [ col for col in X_train.columns if X_train[col].dtypes == 'O']
```

```
categorical
```

```
Out[174...] ['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday']
```

```
In [175...] #display numerical variables
```

```
numerical = [ col for col in X_train.columns if X_train[col].dtypes != 'O']  
numerical
```

```
Out[175...] ['MinTemp',  
            'MaxTemp',  
            'Rainfall',  
            'Evaporation',  
            'Sunshine',  
            'WindGustSpeed',  
            'WindSpeed9am',  
            'WindSpeed3pm',  
            'Humidity9am',  
            'Humidity3pm',  
            'Pressure9am',  
            'Pressure3pm',  
            'Cloud9am',  
            'Cloud3pm',  
            'Temp9am',  
            'Temp3pm',  
            'Year',  
            'Month',  
            'Day']
```

### Engineering missing values in numerical variables

```
In [176...] #check missing values in numerical variables in X_train
```

```
X_train[numerical].isnull().sum()
```

```
Out[176...] MinTemp      495  
            MaxTemp      264  
            Rainfall     1139  
            Evaporation   48718  
            Sunshine     54314  
            WindGustSpeed  7367  
            WindSpeed9am  1086  
            WindSpeed3pm  2094  
            Humidity9am   1449  
            Humidity3pm   2890  
            Pressure9am   11212  
            Pressure3pm   11186  
            Cloud9am     43137  
            Cloud3pm     45768  
            Temp9am      740  
            Temp3pm     2171  
            Year         0  
            Month        0  
            Day          0  
            dtype: int64
```

```
In [177...] #check missing values in numerical variables in X_test
```

```
X_test[numerical].isnull().sum()
```

```
Out[177... MinTemp      142
MaxTemp      58
Rainfall     267
Evaporation  12125
Sunshine     13502
WindGustSpeed 1903
WindSpeed9am 262
WindSpeed3pm 536
Humidity9am   325
Humidity3pm   720
Pressure9am   2802
Pressure3pm   2795
Cloud9am     10520
Cloud3pm     11326
Temp9am      164
Temp3pm      555
Year         0
Month        0
Day          0
dtype: int64
```

```
In [178... #print percentage of missing values in numerical variables in training set
```

```
for col in numerical:
    if X_train[col].isnull().mean() > 0:
        print(col, round(X_train[col].isnull().mean(),4))
```

```
MinTemp 0.0044
MaxTemp 0.0023
Rainfall 0.01
Evaporation 0.4283
Sunshine 0.4775
WindGustSpeed 0.0648
WindSpeed9am 0.0095
WindSpeed3pm 0.0184
Humidity9am 0.0127
Humidity3pm 0.0254
Pressure9am 0.0986
Pressure3pm 0.0983
Cloud9am 0.3792
Cloud3pm 0.4023
Temp9am 0.0065
Temp3pm 0.0191
```

## Assumption

We assume that the data are missing completely at random (MCAR). There are two method which can be used to impute missing values. One is mean or median imputation and other on is random sample imputation. When there are outliers in the dataset, we should use median imputation. So, we will use median imputation because median imputation is robust to outliers.

We will impute missing values with appropriate statistical measures of the data, in this case median. Imputation should be done over the training set, and then propagated to the test set. It means that the statistical measures to be used to fill missing values both train and test set, should be extracted from the train set only. This is to avoid overfitting.

```
In [179... # impute missing values in X_train and X_test with respective column median in X_train
```

```
for df1 in [X_train, X_test] :
    for col in numerical:
        col_median = X_train[col].median()
        df1[col].fillna(col_median, inplace = True)
```

```
In [180... # check again missing values in numerical variables in X_train
```

```
X_train[numerical].isnull().sum()
```

```
Out[180...] MinTemp      0
            MaxTemp      0
            Rainfall      0
            Evaporation    0
            Sunshine      0
            WindGustSpeed  0
            WindSpeed9am   0
            WindSpeed3pm   0
            Humidity9am    0
            Humidity3pm    0
            Pressure9am    0
            Pressure3pm    0
            Cloud9am       0
            Cloud3pm       0
            Temp9am        0
            Temp3pm        0
            Year           0
            Month          0
            Day            0
            dtype: int64
```

```
In [181...] # check missing values in numerical variables in X_test

X_test[numerical].isnull().sum()
```

```
Out[181...] MinTemp      0
            MaxTemp      0
            Rainfall      0
            Evaporation    0
            Sunshine      0
            WindGustSpeed  0
            WindSpeed9am   0
            WindSpeed3pm   0
            Humidity9am    0
            Humidity3pm    0
            Pressure9am    0
            Pressure3pm    0
            Cloud9am       0
            Cloud3pm       0
            Temp9am        0
            Temp3pm        0
            Year           0
            Month          0
            Day            0
            dtype: int64
```

Now, we can see that there are no missing values in numerical columns of training and test set.

## Engineering missing values in categorical variables

```
In [182...] # print percentage of missing values in categorical variables in training set

X_train[categorical].isnull().mean()
```

```
Out[182...] Location      0.000000
            WindGustDir    0.065114
            WindDir9am     0.070134
            WindDir3pm     0.026443
            RainToday      0.010013
            dtype: float64
```

```
In [183...] # print categorical variables with missing data

for col in categorical:
    if X_train[col].isnull().mean() > 0:
        print(col, (X_train[col].isnull().mean()))
```

```
WindGustDir 0.06511419378659213
WindDir9am  0.07013379749283542
WindDir3pm  0.026443026179299188
RainToday   0.01001283471350458
```

```
In [184...] # impute missing categorical variables with most frequent value

for df2 in [X_train, X_test]:
    df2['WindGustDir'].fillna(X_train['WindGustDir'].mode()[0],inplace =True)
    df2['WindDir9am'].fillna(X_train['WindDir9am'].mode()[0],inplace =True)
    df2['WindDir3pm'].fillna(X_train['WindDir3pm'].mode()[0],inplace =True)
    df2['RainToday'].fillna(X_train['RainToday'].mode()[0],inplace =True)
```

```
In [185... # check missing values in categorical variables in X_train
```

```
X_train[categorical].isnull().sum()
```

```
Out[185... Location      0
WindGustDir    0
WindDir9am     0
WindDir3pm     0
RainToday      0
dtype: int64
```

```
In [186... #check missing values in categorical variables in x_test
```

```
X_test[categorical].isnull().sum()
```

```
Out[186... Location      0
WindGustDir    0
WindDir9am     0
WindDir3pm     0
RainToday      0
dtype: int64
```

As a final check, I will check for missing values in X\_train and X\_test.

```
In [187... # check missing values in X_train
```

```
X_train.isnull().sum()
```

```
Out[187... Location      0
MinTemp      0
MaxTemp      0
Rainfall     0
Evaporation  0
Sunshine     0
WindGustDir   0
WindGustSpeed 0
WindDir9am    0
WindDir3pm    0
WindSpeed9am  0
WindSpeed3pm  0
Humidity9am   0
Humidity3pm   0
Pressure9am   0
Pressure3pm   0
Cloud9am      0
Cloud3pm      0
Temp9am       0
Temp3pm       0
RainToday     0
Year          0
Month         0
Day           0
dtype: int64
```

```
In [188... # check missing values in X_test
```

```
X_test.isnull().sum()
```



```
Out[188... Location      0
MinTemp      0
MaxTemp      0
Rainfall     0
Evaporation  0
Sunshine     0
WindGustDir   0
WindGustSpeed 0
WindDir9am   0
WindDir3pm   0
WindSpeed9am 0
WindSpeed3pm 0
Humidity9am   0
Humidity3pm   0
Pressure9am   0
Pressure3pm   0
Cloud9am      0
Cloud3pm      0
Temp9am       0
Temp3pm       0
RainToday     0
Year          0
Month         0
Day           0
dtype: int64
```

We can see that there are no missing values in X\_train and X\_test.

### Engineering outliers in numerical variables

We have seen that the Rainfall, Evaporation, WindSpeed9am and WindSpeed3pm columns contain outliers. We will use top-coding approach to cap maximum values and remove outliers from the above variables.

```
In [189... def max_value(df3, variable, top):
    return np.where(df3[variable]>top, top, df3[variable])

for df3 in [X_train, X_test]:
    df3['Rainfall'] = max_value(df3, 'Rainfall', 3.2)
    df3['Evaporation'] = max_value(df3, 'Evaporation', 21.8)
    df3['WindSpeed9am'] = max_value(df3, 'WindSpeed9am', 55)
    df3['WindSpeed3pm'] = max_value(df3, 'WindSpeed3pm', 57)
```

```
In [190... X_train.Rainfall.max(), X_test.Rainfall.max()
```

```
Out[190... (3.2, 3.2)
```

```
In [191... X_train.Evaporation.max(), X_test.Evaporation.max()
```

```
Out[191... (21.8, 21.8)
```

```
In [192... X_train.WindSpeed9am.max(), X_test.WindSpeed9am.max()
```

```
Out[192... (55.0, 55.0)
```

```
In [193... X_train.WindSpeed3pm.max(), X_test.WindSpeed3pm.max()
```

```
Out[193... (57.0, 57.0)
```

```
In [194... X_train[numerical].describe()
```

Out[194...

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Hu
count	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000
mean	12.193497	23.237216	0.675080	5.151606	8.041154	39.884074	13.978155	18.614756	113754.000000
std	6.388279	7.094149	1.183837	2.823707	2.769480	13.116959	8.806558	8.685862	113754.000000
min	-8.200000	-4.800000	0.000000	0.000000	0.000000	6.000000	0.000000	0.000000	113754.000000
25%	7.600000	18.000000	0.000000	4.000000	8.200000	31.000000	7.000000	13.000000	113754.000000
50%	12.000000	22.600000	0.000000	4.800000	8.500000	39.000000	13.000000	19.000000	113754.000000
75%	16.800000	28.200000	0.600000	5.400000	8.700000	46.000000	19.000000	24.000000	113754.000000
max	33.900000	48.100000	3.200000	21.800000	14.500000	135.000000	55.000000	57.000000	113754.000000



We can now see that the outliers in Rainfall , Evaporation , WindSpeed9am , and WindSpeed3pm columns are all capped.

Encode categorical variables

In [195...

```
categorical
```

Out[195...

```
['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday']
```

In [196...

```
X_train[categorical].head()
```

Out[196...

	Location	WindGustDir	WindDir9am	WindDir3pm	RainToday
110803	Witchcliffe	S	SSE	S	No
87289	Cairns	ENE	SSE	SE	Yes
134949	AliceSprings	E	NE	N	No
85553	Cairns	ESE	SSE	E	No
16110	Newcastle	W	N	SE	No

In [197...

```
!pip install category_encoders
```

Requirement already satisfied: category\_encoders in /usr/local/lib/python3.10/dist-packages (2.6.3)  
Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.10/dist-packages (from category\_encoders) (1.25.2)  
Requirement already satisfied: scikit-learn>=0.20.0 in /usr/local/lib/python3.10/dist-packages (from category\_encoders) (1.2.2)  
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from category\_encoders) (1.11.4)  
Requirement already satisfied: statsmodels>=0.9.0 in /usr/local/lib/python3.10/dist-packages (from category\_encoders) (0.14.2)  
Requirement already satisfied: pandas>=1.0.5 in /usr/local/lib/python3.10/dist-packages (from category\_encoders) (2.0.3)  
Requirement already satisfied: patsy>=0.5.1 in /usr/local/lib/python3.10/dist-packages (from category\_encoders) (0.5.6)  
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.0.5->category\_encoders) (2.8.2)  
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.0.5->category\_encoders) (2023.4)  
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.0.5->category\_encoders) (2024.1)  
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy>=0.5.1->category\_encoders) (1.16.0)  
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.20.0->category\_encoders) (1.4.0)  
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.20.0->category\_encoders) (3.4.0)  
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.10/dist-packages (from statsmodels>=0.9.0->category\_encoders) (24.0)

In [198...

```
#encode RainToday variable

import category_encoders as ce

encoder = ce.BinaryEncoder(cols=['RainToday'])

X_train = encoder.fit_transform(X_train)

X_test = encoder.transform(X_test)
```

In [199...

```
X_train.head()
```

Out[199...

	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	WindDir3pm	...
110803	Witchcliffe	13.9	22.6	0.2	4.8	8.5	S	41.0	SSE	S	...
87289	Cairns	22.4	29.4	2.0	6.0	6.3	ENE	33.0	SSE	SE	...
134949	AliceSprings	9.7	36.2	0.0	11.4	12.3	E	31.0	NE	N	...
85553	Cairns	20.5	30.1	0.0	8.8	11.1	ESE	37.0	SSE	E	...
16110	Newcastle	16.8	29.2	0.0	4.8	8.5	W	39.0	N	SE	...

5 rows × 25 columns



We can see that two additional variables `RainToday_0` and `RainToday_1` are created from `RainToday` variable

Now, We will create the `X_train` training set.

```
In [200... X_train = pd.concat([X_train[numerical], X_train[['RainToday_0', 'RainToday_1']],
                      pd.get_dummies(X_train.Location),
                      pd.get_dummies(X_train.WindGustDir),
                      pd.get_dummies(X_train.WindDir9am),
                      pd.get_dummies(X_train.WindDir3pm)], axis = 1)
```

```
In [201... X_train.head()
```

Out[201...

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity
110803	13.9	22.6	0.2	4.8	8.5	41.0	20.0	28.0	65.0	
87289	22.4	29.4	2.0	6.0	6.3	33.0	7.0	19.0	71.0	
134949	9.7	36.2	0.0	11.4	12.3	31.0	15.0	11.0	6.0	
85553	20.5	30.1	0.0	8.8	11.1	37.0	22.0	19.0	59.0	
16110	16.8	29.2	0.0	4.8	8.5	39.0	0.0	7.0	72.0	

5 rows × 118 columns



Similarly, We will create the `X_test` testing set.

```
In [202... X_test = pd.concat([X_test[numerical], X_test[['RainToday_0', 'RainToday_1']],
                    pd.get_dummies(X_test.Location),
                    pd.get_dummies(X_test.WindGustDir),
                    pd.get_dummies(X_test.WindDir9am),
                    pd.get_dummies(X_test.WindDir3pm)], axis = 1)
```

```
In [203... X_test.head()
```

Out[203...

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity
86232	17.4	29.0	0.0	3.6	11.1	33.0	11.0	19.0	63.0	
57576	6.8	14.4	0.8	0.8	8.5	46.0	17.0	22.0	80.0	
124071	10.1	15.4	3.2	4.8	8.5	31.0	13.0	9.0	70.0	
117955	14.4	33.4	0.0	8.0	11.6	41.0	9.0	17.0	40.0	
133468	6.8	14.3	3.2	0.2	7.3	28.0	15.0	13.0	92.0	

5 rows × 118 columns



We now have training and testing set ready for model building. Before that, we should map all the feature variables onto the same scale. It is called `feature scaling`. We will do it as follows.

## 11. Feature Scaling

```
In [205... X_train.describe()
```

Out [205...

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Hu
count	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000
mean	12.193497	23.237216	0.675080	5.151606	8.041154	39.884074	13.978155	18.614756	65.959641
std	6.388279	7.094149	1.183837	2.823707	2.769480	13.116959	8.806558	8.685862	14.026634
min	-8.200000	-4.800000	0.000000	0.000000	0.000000	6.000000	0.000000	0.000000	1.000000
25%	7.600000	18.000000	0.000000	4.000000	8.200000	31.000000	7.000000	13.000000	52.000000
50%	12.000000	22.600000	0.000000	4.800000	8.500000	39.000000	13.000000	19.000000	62.000000
75%	16.800000	28.200000	0.600000	5.400000	8.700000	46.000000	19.000000	24.000000	70.000000
max	33.900000	48.100000	3.200000	21.800000	14.500000	135.000000	55.000000	57.000000	99.000000

8 rows × 21 columns



In [206...

```
cols = X_train.columns
```

In [207...

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)
```

In [208...

```
X_train = pd.DataFrame(X_train, columns = [cols])
```

In [209...

```
X_test = pd.DataFrame(X_test, columns = [cols])
```


In [210...

```
X_train.describe()
```

Out [210...

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Hu
count	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000	113754.000000
mean	0.484406	0.530004	0.210962	0.236312	0.554562	0.262667	0.254148	0.326575	0.659596
std	0.151741	0.134105	0.369949	0.129528	0.190999	0.101682	0.160119	0.152384	0.140266
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000
25%	0.375297	0.431002	0.000000	0.183486	0.565517	0.193798	0.127273	0.228070	0.520000
50%	0.479810	0.517958	0.000000	0.220183	0.586207	0.255814	0.236364	0.333333	0.620000
75%	0.593824	0.623819	0.187500	0.247706	0.600000	0.310078	0.345455	0.421053	0.700000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	0.990000

8 rows × 118 columns



We now have `X_train` dataset ready to be fed into the Logistic Regression classifier. We will do it as follows.

## 12. Model training

In [211...

```
#train a logistic regression model on the training set
from sklearn.linear_model import LogisticRegression

#instantiate the model
logreg = LogisticRegression(solver='liblinear', random_state=0)

#fit the model
logreg.fit(X_train, y_train)
```

Out [211...

LogisticRegression

LogisticRegression(random\_state=0, solver='liblinear')

## 13. Predict Results

```
In [212... y_pred_test = logreg.predict(X_test)
y_pred_test
```

```
Out[212... array(['No', 'No', 'No', ..., 'No', 'No', 'Yes'], dtype=object)
```

### predict\_proba method

**predict\_proba** method gives the probabilities for the target variable(0 and 1) in this case, in array form.

0 is for probability for no rain and 1 is for probability for rain.

```
In [213... #probability of getting output as 0 - no rain
logreg.predict_proba(X_test)[: ,0]
```

```
Out[213... array([0.91385464, 0.8357268 , 0.82036795, ..., 0.97674976, 0.79859481,
        0.30737669])
```

```
In [214... #probability of getting output as 1 - rain
logreg.predict_proba(X_test)[: ,1]
```

```
Out[214... array([0.08614536, 0.1642732 , 0.17963205, ..., 0.02325024, 0.20140519,
        0.69262331])
```

## 14. Check accuracy score

```
In [215... from sklearn.metrics import accuracy_score
print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test,y_pred_test)))
```

Model accuracy score: 0.8502

Here, **y\_test** are the true class labels and **y\_pred\_test** are the predicted class labels in the test-set.

### Compare the train-set and test-set accuracy

Now, we will compare the train-set and test-set accuracy to check for overfitting.

```
In [216... y_pred_train = logreg.predict(X_train)
y_pred_train
```

```
Out[216... array(['No', 'No', 'No', ..., 'No', 'No', 'No'], dtype=object)
```

```
In [218... print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_pred_train)))
```

Training-set accuracy score: 0.8477

### check for overfitting and underfitting

```
In [219... # print the scores on training and test set
print('Training set score: {:.4f}'.format(logreg.score(X_train,y_train)))
print('Test set score: {:.4f}'.format(logreg.score(X_test, y_test)))
```

Training set score: 0.8477

Test set score: 0.8502

The training-set accuracy score is 0.8477 while the test-set accuracy to be 0.8502. This two values are quite comparable. So, there is no question for overfitting.

In Logistic Regression, we use default value of  $C = 1$ . It provides good performance with approximately 85% accuracy on both the training and the test set. But the model performance on both the training and test set are very comparable. It is likely the case of underfitting.

We will increase  $C$  and fit a more flexible model.

```
In [220... # fit the Logistic Regression model with C=100
#instantiate the model
```

```
logreg100 = LogisticRegression(C=100, solver = 'liblinear', random_state=0)

#fit the model
logreg100.fit(X_train, y_train)
```

Out[220...] **LogisticRegression**  
 LogisticRegression(C=100, random\_state=0, solver='liblinear')

```
In [221...] # print the scores on training and test set

print('Training set score: {:.4f}'.format(logreg100.score(X_train,y_train)))

print('Test set score: {:.4f}'.format(logreg100.score(X_test, y_test)))
```

Training set score: 0.8478  
 Test set score: 0.8505

We can see that, C= 100 results in higher test set accuracy and also a slightly increased training set accuracy. So, we can conclude that a more complex model should perform better.

Now, We will investigate, what happens if we use more regularized model than the default value of C= 1, by setting C= 0.01

```
In [222...] # fit the Logistic Regression model with C=0.01

#instantiate the model

logreg001 = LogisticRegression(C=0.01, solver = 'liblinear', random_state=0)

#fit the model
logreg001.fit(X_train, y_train)
```

Out[222...] **LogisticRegression**  
 LogisticRegression(C=0.01, random\_state=0, solver='liblinear')

```
In [223...] # print the scores on training and test set

print('Training set score: {:.4f}'.format(logreg001.score(X_train,y_train)))

print('Test set score: {:.4f}'.format(logreg001.score(X_test, y_test)))
```

Training set score: 0.8409  
 Test set score: 0.8448

So, if we use more regularized model by setting C=0.01, then both training and test set accuracy decrease relative to the default parameters.

## Compare model accuracy with null accuracy

So, the model accuracy of 0.8501. But, we cannot say that our model is very good based on the above accuracy. We must compare it with the **null accuracy**. Null accuracy is the accuracy that could be achieved by always predicting the most frequent class.

So, we should first check the class distribution in the test set.

```
In [224...] # check class distribution in test set

y_test.value_counts()
```

Out[224...] RainTomorrow  
 No 22067  
 Yes 6372  
 Name: count, dtype: int64

We can see that the occurrences of most frequent class is 22067. So, we can calculate null accuracy by dividing 22067 by total number of occurrences.

```
In [225...] #check null accuracy score

null_accuracy = (22067/(22067+6372))

print('Null accuracy score: {0:0.4f}'.format(null_accuracy))
```

Null accuracy score: 0.7759

We can see that our model accuracy score is 0.8502 but the null accuracy is 0.7759. So, we conclude that our Logistic Regression model is doing a very good job in predicting the class labels.

Now, based on the above analysis we can conclude that our classification model accuracy is very good. Our model is doing a very good job in terms of predicting the class labels.

But, it does not give the underlying distribution of values. Also, it does not tell anything about the type of errors our classifier is making.

We have another tool called **Confusion matrix** that comes to our rescue.

## 15. Confusion Matrix

A confusion matrix is a tool for summarizing the performance of a classification algorithm. A confusion matrix will give us a clear picture of classification model performance and the types of errors produced by the model. It gives us a summary of correct and incorrect predictions broken down by each category. The summary is represented in tabular form.

Four types of outcomes are possible while evaluating a classification model performance. These four outcomes are described below: -

**True Positives (TP)** - True Positives occur when we predict an observation belongs to a certain class and the observation actually belongs to that class.

**True Negatives (TN)** - True Negatives occur when we predict an observation does not belong to a certain class and the observation actually does not belong to that class.

**False Positives (FP)** - False Positives occur when we predict an observation does not belong to a certain class but the observation actually does belong to that class. This type of error is called **Type I error**.

**False Negatives (FN)** - False Negatives occur when we predict an observation does belong to a certain class but the observation actually does not belong to that class. This is a very serious type of error and it is called **Type II error**.

These four outcomes are summarized in a confusion matrix given below.

In [226... *# Print the confusion matrix and slice it into four pieces*

```
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred_test)

print('Confusion Matrix \n\n', cm)

print('\nTrue Positives(TP) = ', cm[0,0])

print('\nTrue Negatives(TN) = ', cm[1,1])

print('\nFalse Positives(FP) = ', cm[0,1])

print('\nFalse Negatives(FN) = ', cm[1,0])
```

Confusion Matrix

```
[[20892  1175]
 [ 3086  3286]]
```

True Positives(TP) = 20892

True Negatives(TN) = 3286

False Positives(FP) = 1175

False Negatives(FN) = 3086

The confusion matrix shows 20892 + 3286 = 24177 correct predictions and 3087 + 1175 = 4262 incorrect predictions.

In this case, we have


- **True Positives** (Actual Positive: 1 and Predict Positive: 1) - 20892
- **True Negatives** (Actual Negative: 0 and Predict Negative: 0) - 3286
- **False Positives** (Actual Negative: 0 but Predict Positive: 1) - 1175 (**Type I error**)
- **False Negatives** (Actual Positive: 1 but Predict Negative: 0) - 3087 (**Type II error**)

In [228... *# visualize confusion matrix with seaborn heatmap*

```
cm_matrix = pd.DataFrame(data = cm, columns = ['Actual Positive:1', 'Actual Negative:0'],
                        index = ['Predict Positive: 1', 'Predict Negative: 0'])
```

```
sns.heatmap(cm_matrix, annot = True, fmt = 'd', cmap='YlGnBu')
```

Out[228... <Axes: >



## 16. Classification Report

### Classification Report

**Classification report** is another way to evaluate the classification model performance. It displays the **precision**, **recall**, **f1** and **support** score for the model. We have described these terms in later.

We can print a classification report as follows:-

```
In [229... from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
No	0.87	0.95	0.91	22067
Yes	0.74	0.52	0.61	6372
accuracy			0.85	28439
macro avg	0.80	0.73	0.76	28439
weighted avg	0.84	0.85	0.84	28439

### Classification accuracy

```
In [231... TP = cm[0,0]
TN = cm[1,1]
FP = cm[0,1]
FN = cm[1,0]
```

```
In [232... # print classification accuracy

classification_accuracy = (TP + TN)/float(TP + TN + FP + FN)

print('Classification accuracy : {0:0.4f}'.format(classification_accuracy))
```

Classification accuracy : 0.8502

### Classification error

```
In [233... # print classification error

classification_error = (FP + FN)/float(TP + TN + FP + FN)

print('Classification error : {0:0.4f}'.format(classification_error))
```



Classification error : 0.1498

## Precision

**Precision** can be defined as the percentage of correctly predicted outcomes out of all the predicted positive outcomes. It can be given as the ratio of true positives (TP) to the sum of the true and false positives (TP + FP)

So, **Precision** identifies the proportion of correctly predicted positive outcome. It is more concerned with the positive class than the negative class

Mathematically, precision can be defined as the ratio of **TP** to **(TP+FP)** .

```
In [234... # print precision score
precision = TP/float(TP +FP)
print('Precision : {0:0.4f}'.format(precision))
```

Precision : 0.9468

## Recall

Recal can be defined as the percentage of correctly predicted positive outcomes out of all the actual positive outcome. It can be given as the ratio of true positives (TP) to the sum of true positives and false negative (TP + FN). **Recall** is also called **Sensitivity**

**Recall** identifies the proportion of correctly predicted actual positives.

Mathematically, Recall can be given as ratio of **TP** to **(TP+FN)** .

```
In [235... recall =TP/ float(TP+FN)
print('Recall or Sensitivity: {0:0.4f}'.format(recall))
```

Recall or Sensitivity: 0.8713

## True Positive Rate

**True Positive Rate** is synonymous with **Recall**.

```
In [236... true_positive_rate = TP/ float(TP + FN)
print('True Positive Rate: {0:0.4f}'.format(true_positive_rate))
```

True Positive Rate: 0.8713

## False Positive Rate

```
In [237... false_positive_rate = FP/ float(FP + TN)
print('False Positive Rate: {0:0.4f}'.format(false_positive_rate))
```

False Positive Rate: 0.2634

## Specificity

```
In [238... specificity = TN/ (TN + FP)
print('Specificity: {0:0.4f}'.format(specificity))
```

Specificity: 0.7366

## f1-score

**f1-score** is weighted harmonic mean of precision and recall. The best possible **f1-score** would be 1.0 and the worst would be 0.0. **f1-score** is the harmonic mean of precision and recall. So, **f1-score** is always lower than accuracy measures as they embed precision and recall into their computation. The weighted average of **f1-score** should be used to compare classifier models, not global accuracy.

## Support

**Support** is the actual number of occurrence of the class in our dataset.

## 17. Adjusting the treshold level

```
In [239... #print the first 10 predicted probabilities of two classes - 0 and 1
```

```
y_pred_prob = logreg.predict_proba(X_test)[0:10]
y_pred_prob
```

```
Out[239... array([[0.91385464, 0.08614536],
       [0.8357268 , 0.1642732 ],
       [0.82036795, 0.17963205],
       [0.99025342, 0.00974658],
       [0.957263 , 0.042737 ],
       [0.97992825, 0.02007175],
       [0.1782761 , 0.8217239 ],
       [0.23474596, 0.76525404],
       [0.90052756, 0.09947244],
       [0.85490676, 0.14509324]])
```

## Observations

- In each row, the numbers sum to 1.
- There are 2 columns which correspond to 2 classes- 0 and 1.
  - Class 0 - predicted probability that there is no rain tomorrow.
  - Class 1 - predicted probability that there is rain tomorrow.
- Importance of predicted probabilities
  - We can rank the observations by probability of rain or no rain.
- predict\_proba process
  - Predicts the probabilities
  - Choose the class with the highest probability
- Classification threshold level
  - There is a classification threshold level of 0.5.
  - Class 1 - probability of rain is predicted if probability > 0.5.
  - Class 0 - probability of no rain is predicted if probability < 0.5.

```
In [240... #store the probabilities in dataframe
```

```
y_pred_prob_df = pd.DataFrame(data = y_pred_prob, columns = ['Prob of - No rain tomorrow (0)', 'Prob of - Rain tomorrow (1)'])
y_pred_prob_df
```

```
Out[240... 
```

	Prob of - No rain tomorrow (0)	Prob of - Rain tomorrow (1)
0	0.913855	0.086145
1	0.835727	0.164273
2	0.820368	0.179632
3	0.990253	0.009747
4	0.957263	0.042737
5	0.979928	0.020072
6	0.178276	0.821724
7	0.234746	0.765254
8	0.900528	0.099472
9	0.854907	0.145093

```
In [241... # print the first 10 predicted probabilities for class 1 - Probability of rain
```

```
logreg.predict_proba(X_test)[0:10, 1]
```

```
Out[241...] array([0.08614536, 0.1642732 , 0.17963205, 0.00974658, 0.042737 ,
      0.02007175, 0.8217239 , 0.76525404, 0.09947244, 0.14509324])
```

```
In [242...] # store the predicted probabilities for class 1 - Probability of rain
y_pred1 = logreg.predict_proba(X_test)[:,-1]
```

```
In [243...] # plot histogram of predicted probabilities

#adjust the font size
plt.rcParams['font.size'] = 12


#plot histogram with 10 bins
plt.hist(y_pred1, bins = 10)

#set the title of predicted probabilities
plt.title('Histogram of predicted probabilities of rain')

#set the x-axis limit
plt.xlim(0,1)

#set the title
plt.xlabel('Predicted probabilities of rain')
plt.ylabel('Frequency')
```

```
Out[243...] Text(0, 0.5, 'Frequency')
```



## Observations

- We can see that the above histogram is highly positive skewed.
- The first column tell us that there are approximately 15000 observations with probability between 0.0 and 0.1
- There are small number of observations with probability > 0.5.
- So, these small number of observations predict that there will be rain tomorrow.
- Majority of observation predict that there will be no rain tomorrow.

## Lower the threshold

```
In [247...] from sklearn.preprocessing import binarize

for i in range(1,5):

    cm1=0

    y_pred1 = logreg.predict_proba(X_test)[:,-1]
    y_pred1 = y_pred1.reshape(-1,1)
```

```

y_pred2 = binarize(y_pred1, threshold = i/10)
y_pred2 = np.where(y_pred2 == 1, 'Yes', 'No')

cm1 = confusion_matrix(y_test, y_pred2)

print('With', i/10, 'threshold the Confusion Matrix is', '\n\n', cm1, '\n\n',
      'with', cm1[0,0]+cm1[1,1], 'correct predictions, ', '\n\n',
      cm1[0,1], 'Type I errors ( False Positives ), ', '\n\n',
      cm1[1,0], 'Type II errors ( False Negatives ), ', '\n\n',
      'Accuracy Score: ', (accuracy_score(y_test, y_pred2)), '\n\n',
      'Sensitivity: ', cm1[1,1]/(float(cm1[1,1]+cm1[1,0])), '\n\n',
      'Specificity: ', cm1[0,0]/(float(cm1[0,0]+cm1[0,1])), '\n\n',
      '=====', '\n\n')

```

With 0.1 threshold the Confusion Matrix is

```
[[12727 9340]
 [ 547 5825]]
```

with 18552 correct predictions,

9340 Type I errors ( False Positives ),

547 Type II errors ( False Negatives ),

Accuracy Score: 0.6523436126446077

Sensitivity: 0.9141556811048337

Specificity: 0.5767435537227534

=====

With 0.2 threshold the Confusion Matrix is

```
[[17066 5001]
 [ 1233 5139]]
```

with 22205 correct predictions,

5001 Type I errors ( False Positives ),

1233 Type II errors ( False Negatives ),

Accuracy Score: 0.7807939800977531

Sensitivity: 0.806497175141243

Specificity: 0.7733720034440568

=====

With 0.3 threshold the Confusion Matrix is

```
[[19079 2988]
 [ 1873 4499]]
```

with 23578 correct predictions,

2988 Type I errors ( False Positives ),

1873 Type II errors ( False Negatives ),

Accuracy Score: 0.829072752206477

Sensitivity: 0.7060577526679221

Specificity: 0.8645941904200843

=====

With 0.4 threshold the Confusion Matrix is

```
[[20192 1875]
 [ 2517 3855]]
```

with 24047 correct predictions,

1875 Type I errors ( False Positives ),

2517 Type II errors ( False Negatives ),

Accuracy Score: 0.8455641900207461

Sensitivity: 0.6049905838041432

Specificity: 0.9150314949925228

=====

## Comments

- In binary problems, the threshold of 0.5 is used by default to convert predicted probabilities into class predictions.
- Threshold can be adjusted to increase sensitivity or specificity.
- Sensitivity and specificity have an inverse relationship. Increase one would always decrease the other and vice versa.
- We can see that increasing the threshold level results in increased accuracy.
- Adjusting the threshold level should be one of the last steps you do in the model-building process.

## 18. ROC - AUC

### ROC Curve

Another tool to measure the classification model performance visually is **ROC Curve**. ROC Curve stands for **Receiver Operating Characteristic Curve**. An **ROC Curve** is a plot which shows the performance of a classification model at various classification threshold levels.

The **ROC Curve** plots the **True Positive Rate (TPR)** against the **False Positive Rate** at various threshold levels.

**True Positive Rate (TPR)** is also called **Recall**. It is defined as the ratio of  $TP$  to  $(TP+FN)$ .

**False Positive Rate** is defined as ratio of  $FP$  to  $(FP +TN)$ .

In the ROC Curve, we will focus on the TPR(True Positive Rate) and FPR(False Positive Rate) of a single point. This will give us the general performance of the ROC Curve which consists of the TPR and FPR at various threshold levels. So, an ROC Curve plots TPR vs FPR at different classification threshold levels. If we lower the threshold levels, it may result in more items being classified as positive. It will increase both True Positive (TP) and False Positive(FP).

```
In [248... # plot ROC Curve

from sklearn.metrics import roc_curve

fpr,tpr, thresholds = roc_curve(y_test,y_pred1, pos_label = 'Yes')

plt.figure(figsize=(6,4))

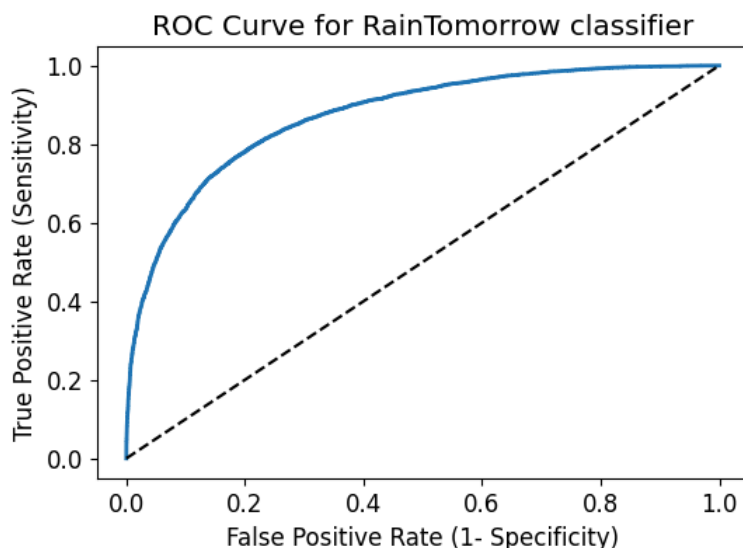
plt.plot(fpr,tpr, linewidth =2)

plt.plot([0,1],[0,1], 'k--' )

plt.rcParams['font.size'] = 12

plt.title('ROC Curve for RainTomorrow classifier')
plt.xlabel('False Positive Rate (1- Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')

plt.show()
```



ROC curve help use to choose a threshold level that balances sensitivity and specificity for a particular context.

## ROC - AUC

**ROC-AUC** stands for **Receiver Operating Characteristic - Area Under Curve**. It is a technique to compare classifier performance. In this technique, we measure the **area under the curve (AUC)**. A perfect classifier will have a ROC - AUC, whereas a purely random classifier will have a ROC AUC equal to 0.5.

So, **ROC AUC** is the percentage of the ROC plot that is underneath the curve.

```
In [251... #compute ROC AUC

from sklearn.metrics import roc_auc_score

ROC_AUC = roc_auc_score(y_test, y_pred1)

print('ROC AUC: {:.4f}'.format(ROC_AUC))
```

ROC AUC: 0.8729

### Comments

- ROC AUC is single number summary of classifier performance. The higher the value, the better the classifier.
- ROC AUC of our model approaches towards 1. So, we can conclude that our classifier does a good job in predicting whether it will rain tomorrow or not.

```
In [252... # calculate cross-validated ROC AUC

from sklearn.model_selection import cross_val_score

Cross_validated_ROC_AUC = cross_val_score(logreg, X_train, y_train, cv = 5, scoring='roc_auc').mean()

print('Cross validated ROC AUC: {:.4f}'.format(Cross_validated_ROC_AUC))
```

Cross validated ROC AUC: 0.8695

## 19.k-Fold Cross Validation

```
In [253... # Applying 5-fold Cross Validation

from sklearn.model_selection import cross_val_score

scores = cross_val_score(logreg, X_train, y_train, cv = 5, scoring='accuracy')

print('Cross-validation scores:{}'.format(scores))
```

Cross-validation scores:[0.84686387 0.84624852 0.84633642 0.84963298 0.84773626]

We can summarize the cross validation accuracy by calculating its mean.

```
In [254... # compute Average cross-validation score

print('Average cross-validation score: {:.4f}'.format(scores.mean()))
```

Average cross-validation score: 0.8474

Our original model score found to be 0.8476. The average cross-validation score is 0.8474. So, we can conclude that cross validation does not result in performance improvement.

## 20. Hyperparameter Optimization using GridSearch CV


```
In [259... from sklearn.model_selection import GridSearchCV

parameters = [{'penalty':['l1', 'l2']],
               {'C':[1, 10, 100, 1000]}]

grid_search = GridSearchCV(estimator = logreg,
                           param_grid = parameters,
                           scoring = 'accuracy',
                           cv = 5,
                           verbose = 0)

grid_search.fit(X_train, y_train)
```

Out[259...



In [260...

```
#examin the best model

#best score achieved during the GridSearchCV
print('GridSeach CV best score : {:.4f}\n\n'.format(grid_search.best_score_))

#print parameters that give the best results
print('Parameters that give the best results :','\n\n',(grid_search.best_params_))

#print estimator that was chosen by GridSearc
print('\n\nEstimate that was chosen by the search :','\n\n',(grid_search.best_estimator_))
```

GridSeach CV best score : 0.8474

Parameters that give the best results :

{'C': 1}

Estimate that was chosen by the search :

LogisticRegression(C=1, random\_state=0, solver='liblinear')

In [261...

```
#calculate GridSearch CV score on test set

print('GridSeach CV score on test set: {0:0.4f}'.format(grid_search.score(X_test, y_test)))
```

GridSeach CV score on test set: 0.8502

## Comments

- Our original model test accuracy is 0.8501 while GridSearch CV accuracy was 0.8502
- We can see that GridSearch CV improve the performance for this particular model.

## 21. Results and conclusion

1. The logistic regression model accuracy score is 0.8501. So, the model does a very good job in predicting whether or not it will rain tomorrow in Australia.
2. Small number of observations predict that there will be rain tomorrow. Majority of observations predict that there will be no rain tomorrow.
3. The model shows no sign of overfitting
4. Increasing the value of C results in higher test set accuracy and also a slightly increased training set accuracy. So, we can conclude that a more complex model should perform better.
5. Increasing the threshold level results in increased accuracy
6. ROC AUC of our model approaches towards 1. So, we can conclude that our classifier does a good job in predicting whether it will rain tomorrow or not.
7. Our Original model accuracy score is 0.8501 whereas accuracy score after RFECV is 0.8500. So, we can obtain approximately similar accuract but with reduced set of features.
8. In the original model, we have FP = 1175 whereas FP1 = 1174. So, we get approximately same number of false positives. Also, FN = 3087 wheareas FN1 = 3091. So, we get slightly higher false negatives.
9. Our, original model score is found to be 0.8476. The average cross validation score is 0.8474. So we can conclude that cross-validation does not result in performance improvement.
10. Our original model test accuracy is 0.8501 while GridSearch CV accuracy is 0.8502. We can see that GridSearch Cv improve the performance for this particular model.

In [ ]:



